

# BizNexus Product Plan: MERN Stack Web App with ML Recommendations

## Functional Specifications by User Role

### Entrepreneurs

- **Profile & Pitch Setup:** Register and create a detailed profile for the startup or project, including business name, industry sector, team members, stage of development, funding needs, and a pitch deck or business plan. Entrepreneurs can upload documents (e.g., executive summary) and tag their startup with relevant categories (e.g., technology, agribusiness).
- **Pitch Ideas & Funding Requests:** Post and publish business pitches or proposals on the platform. Entrepreneurs can specify the amount of funding sought and use forms to structure their requests (problem statement, solution, market, etc.). The system notifies matching investors based on interests.
- **Investor Connections:** Browse and search for investors by criteria (sector interest, investment stage, location). Entrepreneurs can “follow” investors or send private connection requests. A built-in messaging or meeting scheduler lets entrepreneurs initiate conversations with interested investors.
- **Access Resources:** View a knowledge base or resource library (articles, tutorials, templates) relevant to startups and funding. The system recommends resources (articles, guides, courses) based on the entrepreneur’s industry or stage.
- **Progress Tracking & Notifications:** Track the status of each pitch (views by investors, responses received, meetings scheduled). Receive notifications when an investor views or comments on a pitch, or when milestone events occur (e.g. new funding round opportunity).
- **Collaboration:** Optionally invite team members or co-founders. Participate in community forums or Q&A to get feedback on ideas.

### Innovators

- **Idea Sharing:** Register as an innovator and create an “idea profile” summarizing an innovation or concept. Innovators can post idea briefs or prototypes seeking feedback or collaboration.
- **Expert Feedback & Mentorship:** Submit specific questions or issues to the community or a pool of consultants. Receive ratings or comments on ideas. Innovators can apply for mentoring sessions with experts on the platform.

- **Collaboration Tools:** Find and connect with entrepreneurs or other innovators with complementary ideas. Join project teams – the system can list open roles or skill needs (e.g., co-founder with marketing expertise), and innovators can apply.
- **Resource Access:** Similar to entrepreneurs, access tools and templates. The system may recommend incubators or innovation challenges relevant to their idea.
- **Progress Updates:** Keep a log of idea development (e.g., journal, versions). Get notifications on activity (comments, collaboration requests).

## Investors

- **Investor Profile:** Register as an investor, specifying investment criteria: industries/sectors of interest, investment stages (seed, Series A, etc.), geographic focus, and typical ticket size. Optionally link to firm or fund details.
- **Opportunity Browsing:** Search and filter posted pitches by category, stage, region, or funding amount. View a dashboard of recommended startups (see *Machine Learning Integration*). Each pitch view shows key details: summary, team, financials, and attached documents.
- **Pitch Analysis:** For each startup, investors can see analytics (e.g., view count, traction metrics if provided) and leave comments or questions. They can “shortlist” or bookmark pitches. A built-in comment field allows initial feedback or clarification.
- **Investor-Startup Communication:** Upon interest, initiate a private chat or schedule meetings with entrepreneurs. The platform can integrate calendar scheduling or video calls.
- **Investment Tracking:** Maintain a personal portfolio view of past and current investments initiated through the platform. Track statuses (e.g., “in negotiation,” “funded”) and key dates. Investors also receive performance dashboards or alerts on portfolio companies (news updates, milestones).
- **Collaboration:** Join syndicates or co-investment groups: the platform can allow multiple investors to co-fund a deal. Investors can share deal notes or invite other trusted investors to review deals.

## Consultants

- **Consultant Profile:** Register as a consultant or advisor, listing areas of expertise (e.g., finance, legal, marketing), industry experience, credentials, and hourly or project rates. Include sample case studies or links.
- **Service Marketplace:** View a board of “consulting requests” posted by entrepreneurs/innovators seeking help (e.g., “need financial modeling,” “legal advice on incorporation”). Consultants can bid or offer services in response.

- **Direct Offers & Content:** Proactively list services or mini-project offers (e.g., pre-paid packages). Publish articles, toolkits, or webinars to showcase expertise (content discovery). The platform recommends consultants to entrepreneurs based on startup needs (see ML Recs).
- **Engagement Workflow:** When a client (entrepreneur/innovator) wants consulting, the consultant is notified. They can negotiate the scope via chat or schedule a consultation. Optionally, the platform can handle basic contracting and payment (depending on business rules).
- **Feedback & Ratings:** After a consulting engagement, clients can rate and review the consultant. High ratings increase visibility in search/recommendations.

## Admins

- **User Management:** Create, edit, deactivate, or delete user accounts (for any role). Assign roles and permissions. Enforce email verification and profile completeness rules.
- **Content Moderation:** Review and approve content where necessary (e.g. initial pitches or posts for compliance). Moderate community forums and messaging to prevent spam, abuse or inappropriate content. Use automated flags or manual review to handle reports.
- **Platform Analytics:** Access dashboards of key metrics: number of active users by role, number of pitches posted, funding deals completed, site traffic, etc. Use this data to identify growth areas or issues.
- **System Operations:** Manage global settings (e.g. categories, templates, email server). Perform maintenance tasks like database backups, version updates, and security checks. Oversee integration points (e.g., payment gateway if any, ML service).
- **Policy Enforcement:** Ensure users follow platform terms (e.g. no spamming, legitimate business claims). Handle disputes or violations by warning/suspending accounts.
- **Support:** Provide customer support through an admin interface (view support tickets, respond to inquiries).

## Machine Learning Integration

### Recommendation Features

- **Investor Recommendations for Entrepreneurs:** Use collaborative filtering to suggest potential investors to entrepreneurs based on matching criteria and past interactions. For example, if similar startups have attracted certain investors or if an

entrepreneur's tags match an investor's interests, those investors are recommended on the entrepreneur's dashboard. Entrepreneurs see a ranked list of "potential investors" to approach, improving connection speed.

- **Consultant Recommendations for Entrepreneurs/Innovators:** Recommend relevant consultants based on the startup's industry and stated needs. The model analyzes consultant expertise tags and entrepreneur profiles to suggest advisors who have helped similar businesses.
- **Startup Recommendations for Investors:** Recommend promising startups to investors by analyzing patterns of what they have viewed or liked. If an investor frequently looks at tech startups in healthcare, the system highlights new pitches in that intersection. This increases the chance of relevant deal flow.
- **Notifications & Feeds:** The ML system can also power personalized news feeds or alerts. For example, notify an investor when a new startup matching their profile joins, or alert an entrepreneur when a consultant with matching expertise becomes available.

## Data Requirements & Model Integration

- **Data Collection:** Assemble a database of **user profiles** (role, sector tags, past interactions), **item profiles** (startups and investor metadata, consultant expertise), and **interaction data** (which investor viewed or bookmarked which startup, consultant engagements, entrepreneur clicks). Tracking these interactions is critical. As one ML guide notes, "Relevant data, such as product reviews or user ratings, is collected" before building recommendations. Here, equivalent data would be profile views, messages exchanged, and deals executed.
- **Model Type:** Employ a **Collaborative Filtering** approach for many recommendations: it finds similarity between users and items. For example, if Investor A and Investor B have both funded a set of startups in fintech, a new fintech startup that appealed to A might be recommended to B as well. For cold-start scenarios or content-driven matching, use **Content-Based Filtering**: analyze keywords/tags in startup descriptions versus investor interests, or consultant skill tags versus startup problem descriptions. A hybrid model (combining both) can further improve accuracy.
- **Training Pipeline:** Build the recommender as a separate service (e.g., a Python microservice using libraries like scikit-learn or TensorFlow). Periodically (e.g. nightly) train the model on the latest data from MongoDB (user-item matrices, profile features). After training, pre-compute top-N recommendations for each user and store them in the database. The product backend (Node/Express) can fetch these recommended lists via API calls.
- **Integration into MERN:** The MERN backend (Node.js) will call the ML service to update recommendation tables. During user sessions, when an entrepreneur or investor logs in, the React front-end can query endpoints like [/api/recommendations/investors](#) or [/api/recommendations/startups](#),

which return personalized suggestions. This decoupling means the ML component can be implemented in Python but integrated via REST. Alternatively, lightweight on-the-fly scoring (e.g., computing cosine similarities between user and item vectors) could be done in Node if performance allows. The chosen approach ensures recommendations are updated as new users and data arrive.

- **Evaluation:** Monitor key metrics (click-through rates on recommended items, conversion to contacts) to evaluate model performance. A/B testing different algorithms (e.g., matrix factorization vs. nearest-neighbors) can refine the system. Continual retraining ensures the system improves over time.

## Feasibility Analysis

### Technical Feasibility (MERN)

The MERN stack (MongoDB, Express.js, React, Node.js) is well-suited for building a scalable, full-featured web application. It is a **proven, cohesive JavaScript stack**: React provides a dynamic single-page-app front-end, Node/Express serves JSON APIs, and MongoDB handles diverse data (user profiles, pitches, logs). As one source notes, “the MERN stack has emerged as a powerful and popular choice for building scalable applications”. Node.js’s non-blocking, event-driven architecture easily handles many simultaneous users (important for community platforms). Express simplifies API development (authentication, REST endpoints). MongoDB’s flexible document model accommodates evolving data (startups may add fields over time) and supports indexing for fast queries. Overall, the MERN stack allows fast development with a unified language (JavaScript/TypeScript) across client and server.

The integrated machine-learning component can be hosted alongside MERN: for example, a Node.js server can call a Python-based recommendation service or use a Node ML library. Since much of the ML logic is offline (pre-computation), it can run in batch without affecting real-time responsiveness. In short, all required features (real-time updates, media uploads, messaging) are within MERN’s capabilities.

### Scalability & Data Handling

- **Scalability:** The MERN architecture can scale horizontally. For the web app, multiple Node/Express instances can run behind a load balancer (e.g. AWS Elastic Beanstalk or Kubernetes), each stateless for ease of scaling. React assets (JavaScript, CSS) are static files that can be served via CDN, ensuring fast global access. MongoDB can be deployed as a managed cluster (Atlas) with sharding and replica sets to handle high read/write loads and ensure high availability. This supports a growing number of users and data volume.
- **Real-Time Data:** Components like messaging or live notifications can use WebSockets or libraries (Socket.io) on Node, which scale across instances with a message broker (Redis) for pub/sub.
- **Data Storage:** All core data (user accounts, profiles, posts, chats) resides in MongoDB collections. Key indexes should be created on common query fields

(userID, tags, categories). For analytics or logs, consider a separate datastore or a Mongo time-series collection. For heavy-read features (e.g., popular posts), implement caching (e.g., Redis cache) to reduce DB load.

- **Load Testing:** The system should be stress-tested to simulate thousands of concurrent users performing reads and writes (especially the recommender updates). Use Auto-Scaling policies in the cloud to allocate more instances when CPU or request rates spike.
- **Security:** Ensure data is secure with practices like JWT authentication, HTTPS everywhere, and proper input validation. MongoDB user permissions and network restrictions guard against unauthorized access. As the user base grows, regularly audit and patch system components.

## Market Feasibility (Sri Lanka & Global)

Sri Lanka presents a fertile initial market for BizNexus. **SMEs dominate the local economy:** they account for over 75% of registered businesses, about 45% of employment, and 52% of GDP. However, these small and medium enterprises often struggle with financing, networking, and scaling — the very problems BizNexus aims to solve. In recent years, the Sri Lankan startup ecosystem has grown significantly: there are now on the order of **800–900 active startups** supported by government initiatives and accelerators. The country's leadership is actively promoting digital transformation (e.g., a 2030 digital economy strategy) and fostering entrepreneurship through programs (like AccelerateHER). BizNexus aligns well with these trends by connecting entrepreneurs to investors and experts.

The business model also resonates with emerging global trends. Although global venture funding faced a downturn in 2023, demand remains for platforms that streamline deal-sourcing and mentorship. By focusing initially on Sri Lanka's underserved market, BizNexus can capture local momentum before expanding. The solution is easily extensible to similar markets in South Asia or Africa, where SMEs similarly lack connections to capital. In summary, the combination of a large SME base, a vibrant startup scene (hundreds of companies), and supportive policy environments suggests strong market feasibility.

## System Boundaries

### MVP (Minimum Viable Product) Scope:

- **Core Features:** User registration/login (with role selection), profile creation for all roles. Entrepreneurs/innovators: create and post pitches or ideas. Investors: view and search pitches. Consultants: list services or respond to posted consulting requests. Basic messaging/chat between matched users.
- **Basic Matching:** Simple filter/search (by category or keyword) for investors looking at startups, and vice versa. Manual requests to connect (without advanced ML).
- **Admin Panel:** Ability to manage users and moderate content.
- **Exclusions:** No payment processing or handling of actual funds (deals happen off-platform initially). No sophisticated analytics dashboards (just basic metrics). No

mobile app (web-only). Machine learning features (recommendations) may be replaced by simple sorting or personal lists in the MVP.

### Extended Roadmap (Beyond MVP):

- **Machine-Learning Recommendations:** Deploy the recommender engine to suggest investors/consultants/startups, as described above.
- **Advanced Search/Filters:** Enhanced search with faceted filters (e.g. by funding amount, location).
- **Integrated Communication:** Add video conferencing or scheduling integrations for meetings.
- **Transaction Support:** Allow in-platform contracts or escrow for consulting fees or investment (later stage).
- **Analytics & Reporting:** Dashboards for investors (deal funnel analytics) and for admins (user engagement analytics).
- **Mobile App:** Develop native or hybrid mobile versions.
- **Globalization:** Multi-language support (Sinhala, Tamil, English) and local regulations compliance for new regions.
- **Community Features:** Forums, events calendar, or co-innovation challenges.
- **API Ecosystem:** Open APIs for third-party tools (e.g., integration with accounting or CRM systems).

## Business Model

### Revenue Streams

- **Subscription Plans:** Offer tiered subscriptions. For example, entrepreneurs might pay a monthly fee for premium features (featured pitches, detailed investor analytics), while basic accounts remain free. Investors could subscribe for enhanced deal filtering or unlimited messaging (similar to LinkedIn Premium). Consultants might pay a listing fee for exposure. Recurring SaaS revenue from subscriptions provides predictable cash flow.
- **Transaction/Success Fees:** Charge a commission or “success fee” on successful funding deals or consulting contracts initiated through the platform. For example, if an investor closes a funding round with an entrepreneur via BizNexus, the platform takes a small percentage of the deal (1–5%). This aligns the platform’s revenue with user success.

- **Freemium Consulting:** Consultants could use a freemium model – free basic profiles but paid listings or ads for higher visibility. Also, BizNexus could charge for premium consultant services (e.g., a featured placement in the “Top Consultants” list).
- **Advertising & Partnerships:** In later stages, relevant ads (e.g., co-working spaces, business tools) could be served. Partnerships with incubators or accelerators might involve revenue-sharing (e.g. partner pays for leads).
- **Events & Training:** Organize paid webinars or workshops (online/offline) for startups, or sponsor hackathons.

## Go-to-Market Strategy

- **Initial Launch in Sri Lanka:** Focus on major cities (Colombo, Kandy). Work with local entrepreneur organizations (e.g. ICT Agency of Sri Lanka, startup incubators like Hatch, university entrepreneurship cells) to onboard early users.
- **Pilot with Key Users:** Partner with one or two anchor organizations (e.g. a government SME fund or local VC) to provide initial investor and entrepreneur content. This attracts others through credibility.
- **Digital Marketing & Events:** Promote via social media, tech community meetups, and startup events like Global Entrepreneurship Week. Host launch webinars. Leverage success stories (funded deals) as case studies.
- **Strategic Partnerships:** Form alliances with banks (for SME lending programs), legal firms, and educational institutions to integrate their expertise/resources on the platform. These partners can also fund promotional campaigns.
- **Regional Expansion:** After establishing in Sri Lanka, target similar markets in South Asia (e.g., Bangladesh, Nepal) by localizing content and forming local partnerships. Use the proven model as a selling point to international investors interested in emerging markets.

## Scalability Plan

- **Infrastructure:** Host the platform on cloud services (e.g., AWS, Google Cloud) using containerized deployments (Docker, Kubernetes) for easy scaling. Employ managed database and caching services for reliability.
- **Performance Monitoring:** Implement monitoring (e.g., Prometheus, Grafana) to track server load, database performance, and user behavior. Automatically scale resources in response to demand spikes (auto-scaling groups).
- **Data Strategy:** As users grow, periodically archive old logs, ensure indexes are optimized, and possibly shard databases by region or data type. Use CDNs for static



assets and optimize images.

- **Organizational Growth:** Plan to grow the development and support team in phases (e.g., hire additional developers once the user base crosses certain milestones). Adopt agile development with continuous deployment to rapidly iterate based on feedback.
- **Feature Roadmap:** Continuously evaluate user feedback and market trends to prioritize new features (e.g. AI-driven analytics, blockchain for secure contracts). Maintain backward compatibility and data integrity during expansions.
- **Global Compliance:** For international growth, prepare to comply with different data protection laws (e.g. GDPR if reaching EU customers) and financial regulations in each market.
- **Metrics:** Track KPIs such as Monthly Active Users (by role), number of pitches posted, deals closed, and ARPU (average revenue per user). Use these metrics to inform scaling decisions and fundraising rounds.

By following this detailed product plan, BizNexus can evolve from a concept to a robust, scalable MERN-based platform with powerful recommendation features, tailored to the needs of entrepreneurs, innovators, investors, and consultants. Its feasibility is underpinned by strong local market indicators (Sri Lankan SMEs and startups contributing to roughly 52% GDP and over 870 high-tech startups), a proven technical stack, and a clear revenue model aligned with user value. The system's boundaries are well-defined to allow a focused MVP rollout, with a clear roadmap for future growth and expansion.

**Sources:** We base these recommendations on industry-standard practices (e.g., MERN stack capabilities) and market research (e.g., Sri Lankan SME and startup statistics), as well as on best practices for recommendation systems.

## Security Aspects

BizNexus employs a layered security approach across frontend and backend components.

**Frontend (React):** We handle authentication tokens (JWTs) carefully – ideally storing them in secure, *HttpOnly* cookies (with the `Secure` and `SameSite` flags) rather than in `localStorage` to prevent JavaScript access<sup>[owasp.org](https://owasp.org)</sup>. React's default rendering escapes inserted values, which mitigates XSS by design. We **avoid** using `dangerouslySetInnerHTML` with unsanitized content. If we must render raw HTML (e.g. from a WYSIWYG editor), we sanitize it (using libraries like DOMPurify) before use<sup>[stackhawk.com](https://stackhawk.com)</sup><sup>[dmwebsoft.com](https://dmwebsoft.com)</sup>. We also enable HTTPS/TLS for all communications (enforced via HSTS headers with Helmet) so that tokens and data in transit are encrypted<sup>[github.com](https://github.com)</sup><sup>[github.com](https://github.com)</sup>. CSRF attacks are mitigated by using anti-CSRF tokens and SameSite cookies; for example, Express's `csrf` middleware issues a synchronizer token to validate state-changing requests<sup>[dmwebsoft.com](https://dmwebsoft.com)</sup>.

- Frontend XSS/CSRF:** Never interpolate untrusted input into the DOM without sanitization. React's JSX automatically escapes values, but any escape hatch (e.g. `dangerouslySetInnerHTML` or direct DOM manipulation) must be paired with input sanitizers (DOMPurify or similar)[stackhawk.com](https://stackhawk.com). Use `Content-Security-Policy` headers to restrict executable content. Store auth tokens in HttpOnly cookies to prevent scripts from reading them[owasp.org](https://owasp.org); if using local storage, recognize the increased XSS risk. For CSRF protection, either send JWTs in headers (so malicious sites can't automatically send them) or pair cookies with CSRF tokens (and `SameSite=strict`) to ensure requests originate from our app[dmwebsoft.com](https://dmwebsoft.com).
- Backend (Node.js/Express):** The server enforces **Role-Based Access Control (RBAC)** by assigning each user a role (e.g. "admin", "entrepreneur", "investor", "consultant") and protecting routes with middleware. Each request must include a valid JWT (signed with a strong secret and reasonable expiry); middleware (e.g. Passport-JWT or custom) verifies and decodes it. The backend sanitizes and validates all inputs, using libraries like express-validator or Joi. This ensures that query parameters or JSON bodies contain the expected types and formats. Proper input validation is crucial to prevent injection attacks – for example, unsanitized queries against MongoDB can lead to NoSQL injection[dmwebsoft.com](https://dmwebsoft.com). We also set **secure HTTP headers** with Helmet: by default it enables CSP, HSTS (`Strict-Transport-Security: max-age=...`) to prefer HTTPS[github.com](https://github.com), X-Frame-Options to prevent clickjacking, and disables the legacy X-XSS-Protection header to avoid obscure issues. Rate limiting (e.g. via `express-rate-limit` with Redis storage) is applied to all public endpoints to throttle brute-force attempts and mitigate DDoS (limiting requests per IP)[developer.mozilla.org](https://developer.mozilla.org). All responses omit sensitive headers (`Server`, `X-Powered-By`) to not disclose stack info.
- Data Encryption:** All sensitive data is encrypted. Passwords are never stored in plaintext: we use bcrypt with a sufficient work factor (salt rounds) so that hashing is slow enough to resist brute-force and includes a unique salt per password[blog.logrocket.com](https://blog.logrocket.com). For example, using bcrypt's default salt mechanism ensures each stored password hash is unique, preventing rainbow-table attacks[blog.logrocket.com](https://blog.logrocket.com). We also enforce HTTPS/TLS for all client-server communication (HSTS helps ensure browsers use TLS by default[github.com](https://github.com)). For any sensitive fields (e.g. payment info, personal identity), we can use field-level encryption or secure cloud storage; backups and logs containing PII are stored encrypted at rest, and encryption keys are managed securely.
- Compliance (GDPR, Data Protection):** BizNexus collects only essential personal data and gives users control over it. In line with GDPR principles, we minimize data retention (delete obsolete or inactive accounts, purge logs periodically)[securitycompass.com](https://securitycompass.com). We provide clear privacy policies and obtain consent before collecting personal details. Users can request data export or deletion at any time, and we handle these Subject Access Requests (SARs) promptly[cookieyes.com](https://cookieyes.com). All PII (email, profile info) is secured via access controls and encryption[securitycompass.com](https://securitycompass.com). Audit logs track data access and processing activities for accountability. In case of a breach, processes are in place for rapid

response and notification. By design, BizNexus applies “privacy by default/design”: we encrypt or pseudonymize sensitive fields, require strong authentication (e.g. enforce strong passwords, recommend 2FA), and isolate environments.

## Validation Strategy

BizNexus performs comprehensive input validation on both client and server to ensure data quality and security. **Client-side:** React forms provide immediate feedback. We use controlled components and optionally form libraries (e.g. React Hook Form or Formik) to validate required fields, formats, and lengths as the user types. For example, email fields use regex or HTML5 validation to enforce `user@example.com` format, and numeric/length checks prevent out-of-range values. Live validation messages guide the user (e.g. “Invalid email format” or “Password too short”). This enhances UX and catches simple mistakes, but we never trust it for security [formspring.io](https://formspring.io).

- **Client-side Feedback:** Using React’s state, we check inputs on blur or submit. We enforce constraints like minimum/maximum lengths and patterns (e.g. `pattern=/^\S+@\S+\.\S+$/` for emails). Libraries (React Hook Form + Yup, or Formik) simplify schema-based validation and error messages [formspring.io](https://formspring.io). All validation errors are displayed near the field. However, we treat client validation as an aid only.
- **Server-side:** Every request is re-validated on the server using robust middleware. For JSON bodies and parameters, we use `express-validator` or Joi schemas to enforce types and constraints before processing. For example, a pitch-submission endpoint might require `{ title: string(5-100), description: string(0-1000), goalAmount: number(>0) }`. These validators also sanitize inputs (trimming whitespace, escaping HTML, normalizing formats) to remove any malicious content [express-validator.github.io](https://express-validator.github.io). File uploads are checked server-side: we verify file size limits and MIME types, reject non-images, and store uploads on a safe service. Any invalid or malformed input causes a clear HTTP 400 response with a sanitized error message (not exposing internals).
- **Injection Prevention:** As noted, all database queries use parameterized methods. With Mongoose, queries avoid raw string interpolation; we never build queries by concatenating user data. We explicitly whitelist allowed fields. This guards against SQL/NoSQL injection. For instance, using Mongoose’s query APIs (which use server-side BSON translation) mitigates injections by design. We still encode or escape any data rendered in web views.
- **Handling Invalid Data:** On validation failure, the API returns structured error details (e.g. `{ field: "email", message: "Invalid format" }`). The frontend displays these feedbacks to the user. We log validation failures for monitoring (in case of repeated bad input it could signal an attack). Under no circumstances do we let bad data reach business logic or database.

# Database Architecture

We use MongoDB (with Mongoose ODM) to model BizNexus data. Collections are defined for core entities, with clear fields and relationships. All references use ObjectIds. Below is a summary of each collection schema (field: type, notes in *italics*):

yaml

CopyEdit

Users:

```
_id: ObjectId
name: String
email: String (unique, indexed)
passwordHash: String
role: String (enum: ["user", "admin", "consultant", "investor"])
profile: {
  bio: String,
  company: String,
  ...
}
createdAt: Date, updatedAt: Date
```

Pitches:

```
_id: ObjectId
author: ObjectId (ref Users) */ the entrepreneur user who
created this pitch*
title: String
description: String
category: String
goalAmount: Number
currentAmount: Number
attachments: [String] */ file URLs or IDs*
status: String (enum: ["open", "funded", "closed"])
createdAt: Date, updatedAt: Date
```

Investments:

```
_id: ObjectId
investor: ObjectId (ref Users)
pitch: ObjectId (ref Pitches)
amount: Number
investedAt: Date
status: String (enum: ["pending", "completed", "rejected"])
```

Consultants:

```
_id: ObjectId
```

```
user: ObjectId (ref Users)  *// if consultants are also users*
expertise: String
hourlyRate: Number
profile: String  *// detailed text/bio*
createdAt: Date
```

#### Reviews:

```
_id: ObjectId
author: ObjectId (ref Users)
targetType: String (enum: ["pitch", "consultant"])
targetId: ObjectId  *// refers to Pitches or Consultants*
rating: Number (min:1, max:5)
comment: String
createdAt: Date
```

#### Messages:

```
_id: ObjectId
sender: ObjectId (ref Users)
recipient: ObjectId (ref Users)
content: String
sentAt: Date
read: Boolean
```

#### Notifications:

```
_id: ObjectId
user: ObjectId (ref Users)
type: String  *// e.g., "investment", "message", "system"*
payload: Object  *// additional data (e.g., {pitchId, message})*
isRead: Boolean
createdAt: Date
```

#### Logs:

```
_id: ObjectId
user: ObjectId (ref Users, optional)  *// who triggered the
action*
action: String  *// e.g., "LOGIN_SUCCESS", "PITCH_CREATED"*
details: Object  *// any contextual info*
timestamp: Date
```

The above design embeds only truly atomic fields (strings, numbers) and uses **references via ObjectId** for relationships. For example, each `Pitch.author` and `Investment.investor` is an ObjectId pointing to a document in *Users*, establishing a one-to-many relation (one user can author many pitches, invest in many

pitches)[mongodb.com](https://www.mongodb.com). Similarly, `Messages` and `Notifications` tie to users by id. This normalized approach (storing related data in separate collections) improves flexibility and aligns with MongoDB best practices for many-to-many relationships[mongodb.com](https://www.mongodb.com).

Indexes are defined on key fields: e.g. `Users.email` is unique for fast lookups, and common query fields like `Pitches.author`, `Investments.pitch`, and timestamps (`createdAt`) are indexed for performance. The schema ensures data integrity via Mongoose schema types (required fields, enums). Together, this structure captures all BizNexus entities (users, pitches, investments, consultant profiles, reviews, messages, notifications, and audit logs) in a clear, scalable way that maps naturally to MongoDB's document model.